

Username: YALE UNIVERSITY ENGINEERING & APPLIED SCIENCE LIBRARY **Book:** Modern Compiler Implementation in ML. No part of any chapter or book may be reproduced or transmitted in any form by any means without the prior written permission for reprints and excerpts from the publisher of the book or chapter. Redistribution or other use that violates the fair use privilege under U.S. copyright laws (see 17 USC107) or that otherwise violates these Terms of Service is strictly prohibited. Violators will be prosecuted to the full extent of U.S. Federal and Massachusetts laws.

5.2 BINDINGS FOR THE Tiger COMPILER

With what should a symbol table be filled – that is, what is a binding? Tiger has two separate name spaces, one for types and the other for functions and variables. A type identifier will be associated with a `Types.ty`. The `Types` module describes the structure of types, as shown in [Program 5.7](#).

```
structure Types =
struct
  type unique = unit ref

  datatype ty = INT
              | STRING
              | RECORD of (Symbol.symbol * ty) list * unique
              | ARRAY of ty * unique
              | NIL
              | UNIT
              | NAME of Symbol.symbol * ty option ref
end
```

PROGRAM 5.7. Structure `Types`.

The primitive types in Tiger are `int` and `string`; all types are either primitive types or constructed using records and arrays from other (primitive, record, or array) types.

Record types carry additional information: the names and types of the fields. Furthermore, since every “record type expression” creates a new (and different) record type, even if the fields are similar, we have a “unique” value to distinguish it. (The only interesting thing you can do with a `unit ref` is to test it for equality with another one; each `ref` is unique.)

Arrays work just like records: the `ARRAY` constructor carries the type of the array elements, and also a “unique” value to distinguish this array type from all others.

If we were compiling some other language, we might have the following as a legal program:

```
let type a = {x: int, y: int}
    type b = {x: int, y: int}
    var i : a := ...
    var j : b := ...
in i := j
end
```

This is illegal in Tiger, but would be legal in a language where structurally equivalent types are interchangeable. To test type equality in a compiler for such a language, we would need to examine record types field by field, recursively.

However, the following Tiger program is legal, since type `c` is the same as type `a`:

```

let type a = {x: int, y: int}
  type c = a
  var i : a := ...
  var j : c := ...
in i := j
end

```

It is not the type *declaration* that causes a new and distinct type to be made, but the type *expression* {x: int, y:int}.

In Tiger, the expression `nil` belongs to any record type. We handle this exceptional case by inventing a special “nil” type. There are also expressions that return “no value,” so we invent a type `unit`.

When processing mutually recursive types, we will need a place-holder for types whose name we know but whose definition we have not yet seen. The type `NAME(sym, ref(SOME(t)))` is equivalent to type `t`; but `NAME(sym, ref(NONE))` is just the place-holder.

ENVIRONMENTS

The `table` type of the `Symbol` module provides mappings from symbols to bindings. Thus, we will have a *type environment* and a *value environment*. The following Tiger program demonstrates that one environment will not suffice:

```

let type a = int
  var a : a := 5
  var b : a := a
  in b+a
end

```

The symbol `a` denotes the type “a” in syntactic contexts where type identifiers are expected, and the variable “a” in syntactic contexts where variables are expected.

For a type identifier, we need to remember only the type that it stands for. Thus a type environment is a mapping from symbol to `Types.ty` – that is, a `Types.ty Symbol.table`. As shown in [Figure 5.8](#), the `Env` module will contain a `base_tenv` value – the “base” or “predefined” type environment. This maps the symbol `int` to `Ty.INT` and `string` to `Ty.STRING`.

```

signature ENV =
sig
  type access
  type ty
  datatype entry = VarEntry of {ty: ty}
                  | FunEntry of {formals: ty list, result: ty}
  val base_tenv : ty Symbol.table      (* predefined types *)
  val base_venv : entry Symbol.table   (* predefined functions *)
end

```

FIGURE 5.8. Environments for type-checking.

We need to know, for each value identifier, whether it is a variable or a function; if a variable, what is its type; if a function, what are its parameter and result types, and so on. The type `entry` holds all this information, as shown in [Figure 5.8](#); and a value environment is a mapping from symbol to environment-entry.

A variable will map to a `VarEntry` telling its type. When we look up a function we will obtain a `FunEntry` containing:

`formals` The types of the formal parameters.

`result` The type of result returned by the function (or `UNIT`).

For type-checking, only `formals` and `result` are needed; we will add other fields later for translation into intermediate representation.

The `base_venv` environment contains bindings for predefined functions `flush`, `ord`, `chr`, `size`, and so on, described in [Appendix A](#).

Environments are used during the type-checking phase.

As types, variables, and functions are declared, the type-checker augments the environments; they are consulted for each identifier that is found during processing of expressions (type-checking, intermediate code generation).

